

⌕ B I <> 🔗 📄 🔊 ⚙️ ⚙️ — 📄 ☺️ ⋮

```
*Problem Statement - Implement data processing tasks using Spark
**
*** github link :
```

Problem Statement - Implement data processing tasks using Spark  
DataFrame API.\*\*\* github link :

```
import pandas as pd
import matplotlib.pyplot as plt
```

## ✓ 1. Utilize PySpark DataFrame API for tasks like filtering, grouping, and aggregation.

### 1. \*Installing PySpark\*:

```
!pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
```

### 2. Creating a Spark Session:

```
from pyspark.sql import SparkSession

# Create a SparkSession
spark = SparkSession.builder.appName("SparkSession").getOrCreate()
```

### 3. Loading Dataset:

```
# Import necessary libraries
from pyspark.sql.functions import col

# Load the dataset into a DataFrame
file_path = "bank.csv"
df = spark.read.csv(file_path, header=True, inferSchema=True)

# Show the schema of the DataFrame
df.printSchema()
```

```
root
 |-- age: integer (nullable = true)
 |-- job: string (nullable = true)
 |-- marital: string (nullable = true)
 |-- education: string (nullable = true)
 |-- default: string (nullable = true)
 |-- balance: integer (nullable = true)
 |-- housing: string (nullable = true)
 |-- loan: string (nullable = true)
 |-- contact: string (nullable = true)
 |-- day: integer (nullable = true)
 |-- month: string (nullable = true)
 |-- duration: integer (nullable = true)
 |-- campaign: integer (nullable = true)
 |-- pdays: integer (nullable = true)
 |-- previous: integer (nullable = true)
 |-- poutcome: string (nullable = true)
 |-- y: string (nullable = true)
```

```
# Display the first few rows of the DataFrame
df.show()
```

```
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|age|          job|marital|education|default|balance|housing|loan|  contact|day|month|duration|campaign|pdays|previous|po
```

30	unemployed	married	primary	no	1787	no	no	cellular	19	oct	79	1	-1	0	u
33	services	married	secondary	no	4789	yes	yes	cellular	11	may	220	1	339	4	f
35	management	single	tertiary	no	1350	yes	no	cellular	16	apr	185	1	330	1	f
30	management	married	tertiary	no	1476	yes	yes	unknown	3	jun	199	4	-1	0	u
59	blue-collar	married	secondary	no	0	yes	no	unknown	5	may	226	1	-1	0	u
35	management	single	tertiary	no	747	no	no	cellular	23	feb	141	2	176	3	f
36	self-employed	married	tertiary	no	307	yes	no	cellular	14	may	341	1	330	2	
39	technician	married	secondary	no	147	yes	no	cellular	6	may	151	2	-1	0	u
41	entrepreneur	married	tertiary	no	221	yes	no	unknown	14	may	57	2	-1	0	u
43	services	married	primary	no	-88	yes	yes	cellular	17	apr	313	1	147	2	f
39	services	married	secondary	no	9374	yes	no	unknown	20	may	273	1	-1	0	u

43	admin.	married	secondary	no	264	yes	no	cellular	17	apr	113	2	-1	0	u
36	technician	married	tertiary	no	1109	no	no	cellular	13	aug	328	2	-1	0	u
20	student	single	secondary	no	502	no	no	cellular	30	apr	261	1	-1	0	u
31	blue-collar	married	secondary	no	360	yes	yes	cellular	29	jan	89	1	241	1	f
40	management	married	tertiary	no	194	no	yes	cellular	29	aug	189	2	-1	0	u
56	technician	married	secondary	no	4073	no	no	cellular	27	aug	239	5	-1	0	u
37	admin.	single	tertiary	no	2317	yes	no	cellular	20	apr	114	1	152	2	f
25	blue-collar	single	primary	no	-221	yes	no	unknown	23	may	250	1	-1	0	u
31	services	married	secondary	no	132	no	no	cellular	7	jul	148	1	152	1	

only showing top 20 rows

## Task 1: Filtering

```
import matplotlib.pyplot as plt
```

```
# Task 1: Filtering
```

```
# Filter the DataFrame to select rows where the 'Education' column is 'primary'
```

```
primary_df = df.filter(col("education") == "primary")
```

```
primary_df.show()
```

```
# Convert the filtered DataFrame to Pandas
```

```
primary_pandas_df = primary_df.toPandas()
```

```
→
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous
	30	unemployed	married	primary	no	1787	no	no	cellular	19	oct	79	1	-1	0
	43	services	married	primary	no	-88	yes	yes	cellular	17	apr	313	1	147	2
	25	blue-collar	single	primary	no	-221	yes	no	unknown	23	may	250	1	-1	0
	55	blue-collar	married	primary	no	627	yes	no	unknown	5	may	247	1	-1	0
	78	retired	divorced	primary	no	229	no	no	telephone	22	oct	97	1	-1	0
	55	blue-collar	married	primary	no	145	no	no	telephone	2	feb	59	3	5	2
	26	blue-collar	married	primary	no	0	yes	no	unknown	21	may	425	1	-1	0
	32	entrepreneur	single	primary	yes	-849	yes	yes	cellular	4	feb	204	1	-1	0
	45	blue-collar	divorced	primary	no	844	no	no	unknown	5	jun	1018	3	-1	0
	41	blue-collar	married	primary	no	-516	no	yes	telephone	8	jul	554	3	-1	0
	50	blue-collar	divorced	primary	no	388	no	no	cellular	5	feb	701	1	-1	0
	60	retired	married	primary	no	5	no	no	cellular	26	aug	63	2	-1	0
	51	blue-collar	married	primary	no	1466	yes	no	unknown	7	may	406	2	-1	0
	35	blue-collar	single	primary	no	293	yes	no	unknown	30	may	521	2	-1	0
	39	blue-collar	married	primary	no	111	no	no	cellular	18	nov	201	2	-1	0
	34	blue-collar	married	primary	no	455	yes	no	unknown	20	jun	372	3	-1	0
	31	unemployed	single	primary	no	406	no	no	cellular	4	feb	736	1	-1	0
	57	services	single	primary	no	3777	yes	no	telephone	13	may	65	2	-1	0
	56	unemployed	married	primary	no	3391	no	no	cellular	21	apr	243	1	-1	0
	27	housemaid	married	primary	no	0	yes	no	cellular	23	jul	435	3	-1	0

only showing top 20 rows

## Task 2: Grouping and Aggregation

```
# Task 2: Grouping and Aggregation
```

```
# Group the DataFrame by 'marital' and 'job', then calculate the average the 'duration' for each group
```

```
grouped_df = df.groupBy("marital", "job").agg({"duration": "sum"})
```

```
grouped_df.show()
```

```
# Convert the grouped DataFrame to Pandas
```

```
grouped_pandas_df = grouped_df.toPandas()
```

```
# Pivot the DataFrame for easy plotting
```

```
pivot_df = grouped_pandas_df.pivot(index='marital', columns='job', values='sum(duration)')
```

```
# Plot a grouped bar chart
```

```
ax = pivot_df.plot(kind='bar', figsize=(12, 8), rot=45)
```

```
ax.set_xlabel('month')
```

```
ax.set_ylabel('Sum of Campaign')
```

```
ax.set_title('Campaign Count Grouped by housing loan and Month')
```


```
plt.legend(title='Month')
```

```
plt.show()
```

```

+-----+-----+-----+
| marital|      job|sum(duration)|
+-----+-----+-----+
| married|self-employed|      29055|
| divorced|management|      30232|
| divorced|self-employed|      3880|
|  single|self-employed|      15400|
| married|entrepreneur|      35516|
| divorced|      unknown|       385|
| married|      admin.|      58628|
| divorced|      admin.|      19885|
| married|unemployed|      24440|
|  single|management|      70045|
| married|services|      63662|
|  single|services|      29541|
| married|blue-collar|     178968|
| divorced|services|      16254|
| married|technician|      95281|
| divorced|entrepreneur|      6306|
|  single|housemaid|       5081|
|  single|retired|       3328|
|  single|      admin.|      33659|
| divorced|technician|      20328|
+-----+-----+-----+
only showing top 20 rows

```



### 3. Handle missing data and outliers using Spark's data preprocessing capabilities.

#### 1. Fill Missing Values:

For numeric columns, you can fill missing values with mean, median, or any other statistical measure.

```
df.show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| age|      job|marital|education|default|balance|housing|loan|contact|day|month|duration|campaign|pdays|previous|po
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 30|unemployed|married|primary|no|1787|no|no|cellular|19|oct|79|1|-1|0|u
| 33|services|married|secondary|no|4789|yes|yes|cellular|11|may|220|1|339|4|f
| 35|management|single|tertiary|no|1350|yes|no|cellular|16|apr|185|1|330|1|f
| 30|management|married|tertiary|no|1476|yes|yes|unknown|3|jun|199|4|-1|0|u

```

```

59| blue-collar|married|secondary|no|0|yes|no|unknown|5|may|226|1|-1|0|u
35| management|single|tertiary|no|747|no|no|cellular|23|feb|141|2|176|3|f
36|self-employed|married|tertiary|no|307|yes|no|cellular|14|may|341|1|330|2|
39| technician|married|secondary|no|147|yes|no|cellular|6|may|151|2|-1|0|u
41| entrepreneur|married|tertiary|no|221|yes|no|unknown|14|may|57|2|-1|0|u
43| services|married|primary|no|-88|yes|yes|cellular|17|apr|313|1|147|2|f
39| services|married|secondary|no|9374|yes|no|unknown|20|may|273|1|-1|0|u
43| admin.|married|secondary|no|264|yes|no|cellular|17|apr|113|2|-1|0|u
36| technician|married|tertiary|no|1109|no|no|cellular|13|aug|328|2|-1|0|u
20| student|single|secondary|no|502|no|no|cellular|30|apr|261|1|-1|0|u
31| blue-collar|married|secondary|no|360|yes|yes|cellular|29|jan|89|1|241|1|f
40| management|married|tertiary|no|194|no|yes|cellular|29|aug|189|2|-1|0|u
56| technician|married|secondary|no|4073|no|no|cellular|27|aug|239|5|-1|0|u
37| admin.|single|tertiary|no|2317|yes|no|cellular|20|apr|114|1|152|2|f
25| blue-collar|single|primary|no|-221|yes|no|unknown|23|may|250|1|-1|0|u
31| services|married|secondary|no|132|no|no|cellular|7|jul|148|1|152|1|

```

only showing top 20 rows

```

# find missing values using pyspark
from pyspark.sql.functions import col,isnan, when, count

df.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in df.columns]
).show()

```

```

+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|age|job|marital|education|default|balance|housing|loan|contact|day|month|duration|campaign|pdays|previous|poutcome| y|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|0|
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

## ✓ Implement a Classification AI Model using Spark

```

# import pyspark ml
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

```

```

# train test split
train_df, test_df = df.randomSplit([0.7, 0.3], seed=42)
# train the random forest classifier

```

## ✓ Task

Use spark ml randomforestclassifier to train the model with train\_df and test\_df. Explain the selected code. Explain the error in the selected code. If possible, fix the error and incorporate the changes into the existing code. Otherwise, try to diagnose the error.

## ✓ Identify categorical columns

Subtask:

Determine which columns in the DataFrame are of string type and need to be converted.

**Reasoning:** Iterate through the columns and check their data types to identify string columns.

```

string_cols = [col_name for col_name, col_type in df.dtypes if col_type == 'string']
print("String columns:", string_cols)

```

```

String columns: ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome', 'y']

```

## ✓ Apply string indexing

Subtask:

Use StringIndexer to convert each categorical string column into a numerical index column.

**Reasoning:** Convert the categorical string columns to numerical indices using StringIndexer as identified in the previous step.

```
from pyspark.ml.feature import StringIndexer

string_cols = [col_name for col_name, col_type in df.dtypes if col_type == 'string']
indexed_cols = [col_name + "_indexed" for col_name in string_cols]

indexer = StringIndexer(inputCols=string_cols, outputCols=indexed_cols)
indexed_df = indexer.fit(df).transform(df)

indexed_df.printSchema()
```

```
↳ root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- balance: integer (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- day: integer (nullable = true)
|-- month: string (nullable = true)
|-- duration: integer (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- y: string (nullable = true)
|-- job_indexed: double (nullable = false)
|-- marital_indexed: double (nullable = false)
|-- education_indexed: double (nullable = false)
|-- default_indexed: double (nullable = false)
|-- housing_indexed: double (nullable = false)
|-- loan_indexed: double (nullable = false)
|-- contact_indexed: double (nullable = false)
|-- month_indexed: double (nullable = false)
|-- poutcome_indexed: double (nullable = false)
|-- y_indexed: double (nullable = false)
```

## ✓ Apply one-hot encoding

Subtask:

Use OneHotEncoder to convert the indexed columns into one-hot encoded vector columns.

**Reasoning:** Import OneHotEncoder and apply it to the indexed columns to create one-hot encoded vector columns as per the instructions.

```
from pyspark.ml.feature import OneHotEncoder

indexed_cols = [col_name + "_indexed" for col_name, col_type in df.dtypes if col_type == 'string']
encoded_cols = [col_name + "_vec" for col_name in indexed_cols]

encoder = OneHotEncoder(inputCols=indexed_cols, outputCols=encoded_cols)
encoded_df = encoder.fit(indexed_df).transform(indexed_df)

encoded_df.printSchema()
```

```
↳ root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- balance: integer (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- day: integer (nullable = true)
|-- month: string (nullable = true)
|-- duration: integer (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- y: string (nullable = true)
|-- job_indexed: double (nullable = false)
|-- marital_indexed: double (nullable = false)
|-- education_indexed: double (nullable = false)
|-- default_indexed: double (nullable = false)
```

```

|-- housing_indexed: double (nullable = false)
|-- loan_indexed: double (nullable = false)
|-- contact_indexed: double (nullable = false)
|-- month_indexed: double (nullable = false)
|-- poutcome_indexed: double (nullable = false)
|-- y_indexed: double (nullable = false)
|-- job_indexed_vec: vector (nullable = true)
|-- marital_indexed_vec: vector (nullable = true)
|-- education_indexed_vec: vector (nullable = true)
|-- default_indexed_vec: vector (nullable = true)
|-- housing_indexed_vec: vector (nullable = true)
|-- loan_indexed_vec: vector (nullable = true)
|-- contact_indexed_vec: vector (nullable = true)
|-- month_indexed_vec: vector (nullable = true)
|-- poutcome_indexed_vec: vector (nullable = true)
|-- y_indexed_vec: vector (nullable = true)

```

## ✓ Update vectorassembler

### Subtask:

Include the newly created one-hot encoded columns in the VectorAssembler input columns.

**Reasoning:** Identify the numerical and one-hot encoded columns and update the VectorAssembler.

```

# Apply the assembler transformation to the encoded_df
assembled_df = assembler.transform(encoded_df)

# Split the assembled_df into training and testing sets
train_data, test_data = assembled_df.randomSplit([0.7, 0.3], seed=42)

train_data.printSchema()
test_data.printSchema()

```

```

↔ root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- balance: integer (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- day: integer (nullable = true)
|-- month: string (nullable = true)
|-- duration: integer (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- y: string (nullable = true)
|-- job_indexed: double (nullable = false)
|-- marital_indexed: double (nullable = false)
|-- education_indexed: double (nullable = false)
|-- default_indexed: double (nullable = false)
|-- housing_indexed: double (nullable = false)
|-- loan_indexed: double (nullable = false)
|-- contact_indexed: double (nullable = false)
|-- month_indexed: double (nullable = false)
|-- poutcome_indexed: double (nullable = false)
|-- y_indexed: double (nullable = false)
|-- job_indexed_vec: vector (nullable = true)
|-- marital_indexed_vec: vector (nullable = true)
|-- education_indexed_vec: vector (nullable = true)
|-- default_indexed_vec: vector (nullable = true)
|-- housing_indexed_vec: vector (nullable = true)
|-- loan_indexed_vec: vector (nullable = true)
|-- contact_indexed_vec: vector (nullable = true)
|-- month_indexed_vec: vector (nullable = true)
|-- poutcome_indexed_vec: vector (nullable = true)
|-- y_indexed_vec: vector (nullable = true)
|-- features: vector (nullable = true)

root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- balance: integer (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)

```

```
-- contact: string (nullable = true)
-- day: integer (nullable = true)
-- month: string (nullable = true)
-- duration: integer (nullable = true)
-- campaign: integer (nullable = true)
-- pdays: integer (nullable = true)
-- previous: integer (nullable = true)
-- poutcome: string (nullable = true)
-- y: string (nullable = true)
```

## ✓ Train the model

### Subtask:

Retrain the RandomForestClassifier with the transformed data.

**Reasoning:** Retrain the RandomForestClassifier model with the transformed data.

```
# Initialize and train the RandomForestClassifier model
# Use 'y_indexed' as the label column as it's the numerical representation of the target variable
rf_model = RandomForestClassifier(labelCol="y_indexed", featuresCol="features", numTrees=10)
rf_model = rf_model.fit(train_data)

# Make predictions on the test data
predictions = rf_model.transform(test_data)

# Evaluate the model
# Use 'y_indexed' as the label column for evaluation
evaluator = MulticlassClassificationEvaluator(labelCol="y_indexed", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)

print(f"Test Accuracy = {accuracy}")
```

🔄 Test Accuracy = 0.9968944099378882

### Summary:

#### Data Analysis Key Findings

- The categorical string columns in the dataset were identified as: `job`, `marital`, `education`, `default`, `housing`, `loan`, `contact`, `month`, `poutcome`, and `y`.
- These string columns were successfully converted into numerical indexed columns using `StringIndexer`, adding columns with a `_indexed` suffix (e.g., `job_indexed`).
- The indexed columns were further transformed into one-hot encoded vector columns using `OneHotEncoder`, resulting in columns with a `_vec` suffix (e.g., `job_indexed_vec`).
- The `VectorAssembler` was updated to include both the numerical columns and the newly created one-hot encoded vector columns as features.
- The data was successfully split into training and testing sets *after* applying the `VectorAssembler` transformation to the `DataFrame` containing the one-hot encoded features.
- A `RandomForestClassifier` model was initialized with `y_indexed` as the label column and `features` as the features column.
- The model was successfully trained on the `train_data` and used to make predictions on the `test_data`.
- The model's performance was evaluated using `MulticlassClassificationEvaluator`, yielding a test accuracy of approximately 0.997.

#### Insights or Next Steps

- The high accuracy suggests the model is performing very well on the test set; further investigation into potential overfitting or evaluation with different metrics (like precision, recall, F1-score) on a validation set would be beneficial.
- Consider exploring hyperparameter tuning for the `RandomForestClassifier` (e.g., `numTrees`, `maxDepth`, `maxBins`) to potentially improve performance further or address overfitting.